


Announcements

- Course TA: Danny Kumar (ndkumar@cs.unc.edu)
- Guest lectures next week
- Assignment 1 out, due Tuesday, September 14
- Course discussion board
 - You should be able to find the discussion thread “AI in the news/on the Web” on Blackboard
 - For participation points: at least two posts during the semester


Solving problems by searching

Chapter 3



Search

- We will consider the problem of designing **goal-based agents** in **observable, deterministic, discrete, known** environments
- Example:




Search

- We will consider the problem of designing **goal-based agents** in **observable, deterministic, discrete, known** environments
 - The solution is a fixed sequence of actions
 - Search is the process of looking for the sequence of actions that reaches the goal
 - Once the agent begins executing the search solution, it can ignore its percepts (**open-loop system**)

Search problem components

- **Initial state**
- **Actions**
- **Transition model**
 - What is the result of performing a given action in a given state?
- **Goal state**
- **Path cost**
 - Assume that it is a sum of nonnegative *step costs*



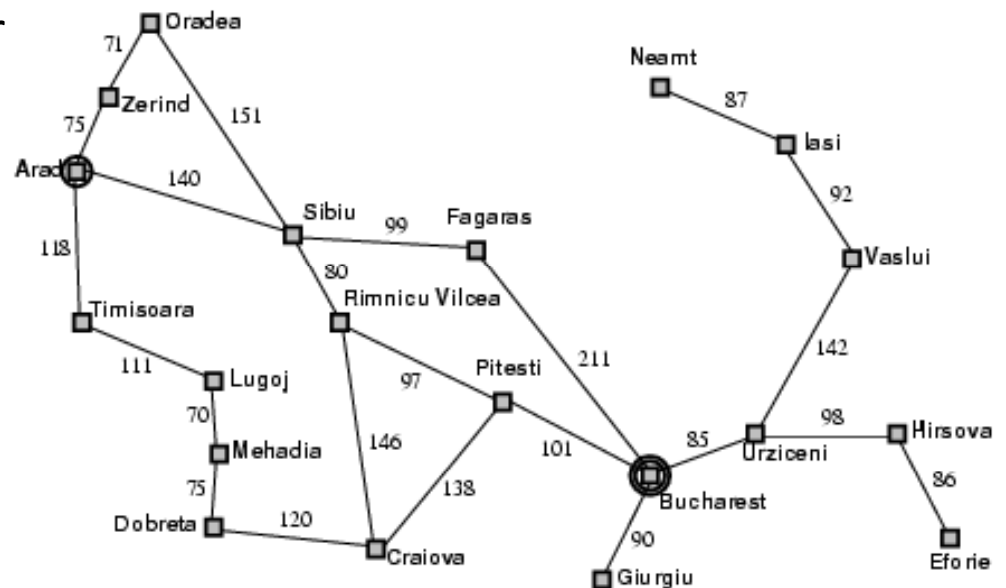
- The **optimal solution** is the sequence of actions that gives the lowest path cost for reaching the goal

Example: Romania

- On vacation in Romania; currently in Arad
- Flight leaves tomorrow from Bucharest




- **Initial state**
 - Arad
- **Actions**
 - Go from one city to another
- **Transition model**
 - If you go from city A to city B, you end up in city B
- **Goal state**
 - Bucharest
- **Path cost**
 - Sum of edge costs



State space


- The initial state, actions, and transition model define the **state space** of the problem
 - The set of all states reachable from initial state by any sequence of actions
 - Can be represented as a **directed graph** where the nodes are states and links between nodes are actions
- What is the state space for the Romania problem?

Example: Vacuum world



- **States**
 - Agent location and dirt location
 - How many possible states?
 - What if there are n possible locations?
- **Actions**
 - Left, right, suck
- **Transition model**

Vacuum world state space graph



Example: The 8-puzzle

- **States**

- Locations of tiles

- 8-puzzle: 181,440 states
 - 15-puzzle: 1.3 trillion states
 - 24-puzzle: 10^{25} states

7	2	4
5		6
8	3	1

Start State

- **Actions**

- Move blank left, right, up, down

- **Path cost**


- 1 per move

	1	2
3	4	5
6	7	8

Goal State

- Optimal solution of n-Puzzle is NP-hard

Example: Robot motion planning



- **States**
 - Real-valued coordinates of robot joint angles
- **Actions**
 - Continuous motions of robot joints
- **Goal state**
 - Desired final configuration (e.g., object is grasped)
- **Path cost**
 - Time to execute, smoothness of path, etc.

Other Real-World Examples

- Routing
- Touring
- VLSI layout
- Assembly sequencing
- Protein design

Search


- Given:
 - Initial state
 - Actions
 - Transition model
 - Goal state
 - Path cost
- How do we find the optimal solution?
 - How about building the state space and then using Dijkstra's shortest path algorithm?
 - The state space is huge!
 - Complexity of Dijkstra's is $O(E + V \log V)$, where V is the size of the state space

Tree Search

- Let's begin at the start node and **expand** it by making a list of all possible successor states
- Maintain a **fringe** or a list of unexpanded states
- At each step, pick a state from the fringe to expand
- Keep going until you reach the goal state
- Try to expand as few states as possible

Search tree

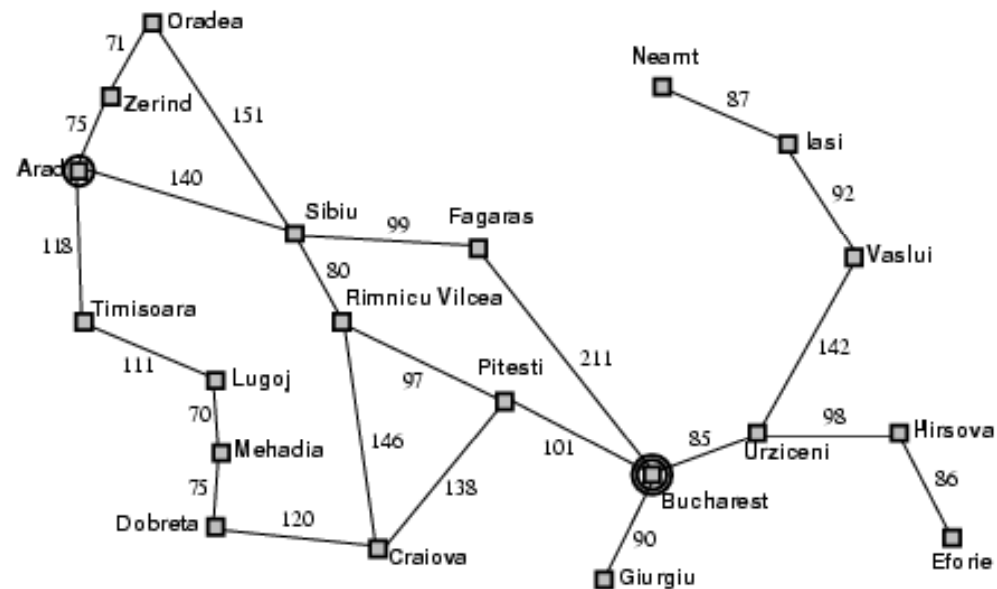
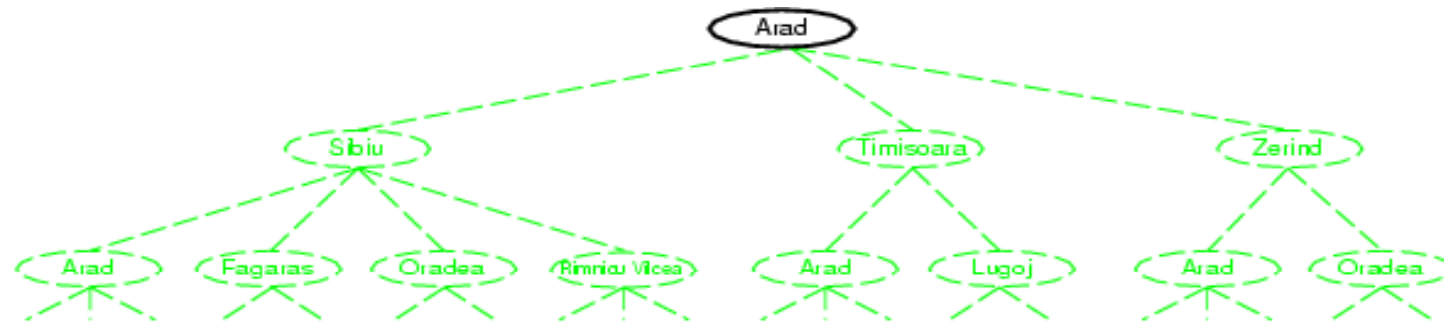
- “What if” tree of possible actions and outcomes
- The root node corresponds to the starting state
- The children of a node correspond to the **successor states** of that node's state
- A path through the tree corresponds to a sequence of actions
 - A solution is a path ending in the goal state
- Nodes vs. states
 - A state is a representation of a physical configuration, while a node is a data structure that is part of the search tree




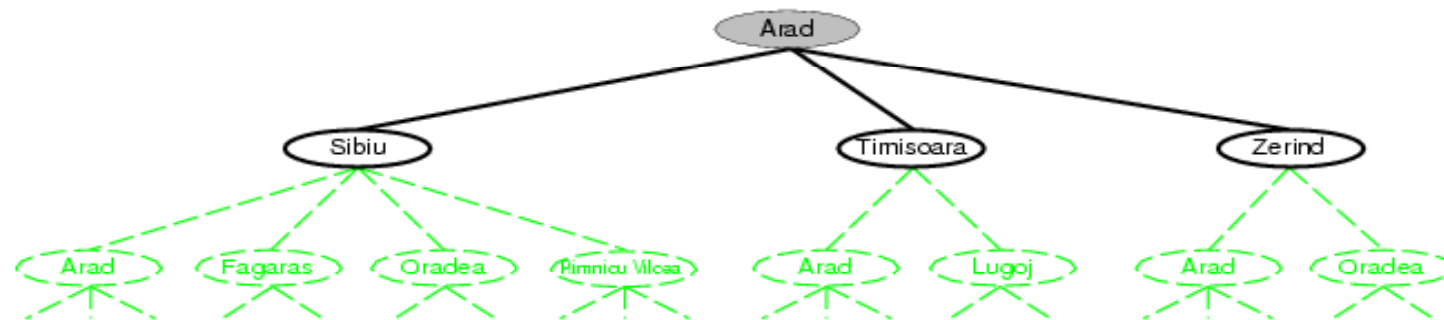
Tree Search Algorithm Outline

- Initialize the **fringe** using the **starting state**
- While the fringe is not empty
 - Choose a fringe node to expand according to **search strategy**
 - If the node contains the **goal state**, return solution
 - Else **expand** the node and add its children to the fringe


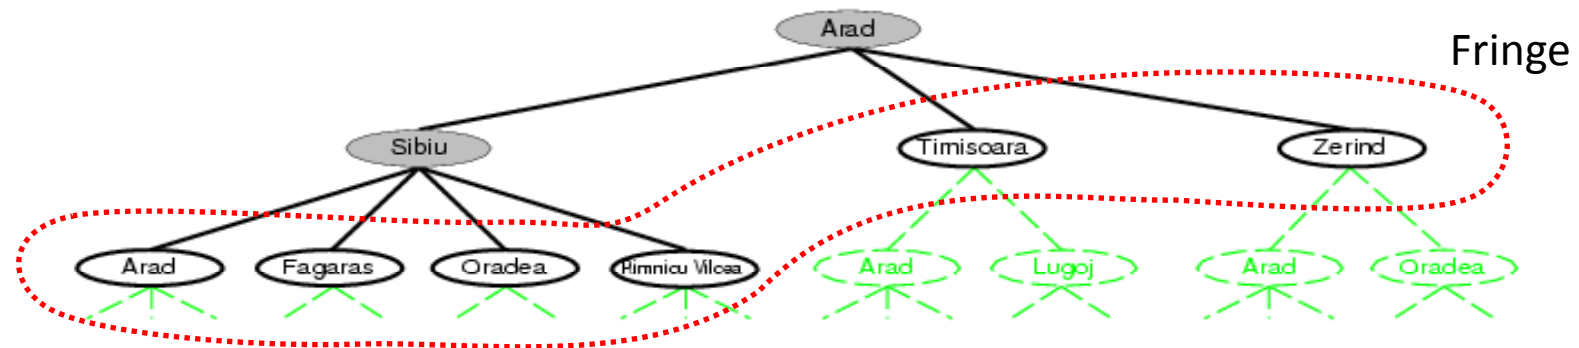
Tree search example



Tree search example



Tree search example



Search strategies

- A **search strategy** is defined by picking the order of node expansion
- Strategies are evaluated along the following dimensions:
 - **Completeness**: does it always find a solution if one exists?
 - **Optimality**: does it always find a least-cost solution?
 - **Time complexity**: number of nodes generated
 - **Space complexity**: maximum number of nodes in memory
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum length of any path in the state space (may be infinite)

Uninformed search strategies

- **Uninformed** search strategies use only the information available in the problem definition
- Breadth-first search
- Uniform-cost search
- Depth-first search
- Iterative deepening search